

Free Energy Calculation Methods in Biomolecular Systems: Theory and Practice

Jaeoh Shin
Korea Institute for Advanced Study
Seoul, Korea

January 8, 2026

Version 1.0

Abstract

Free energy calculations are fundamental to understanding molecular recognition, protein stability, and drug binding in biomolecular systems. This document provides a comprehensive introduction to the three major computational approaches: Free Energy Perturbation (FEP), Thermodynamic Integration (TI), and nonequilibrium methods based on the Jarzynski equality. We emphasize fundamental concepts and their connection to practical implementation, bridging rigorous statistical mechanics with modern molecular dynamics software. The material is designed for researchers entering the field of computational drug discovery who want to understand both the theoretical foundations and practical implementation of free energy methods, and includes a hands-on tutorial using GROMACS and pmx. We illustrate key concepts using a nonequilibrium switching approach applied to CDK8 kinase inhibitors—a relevant drug discovery application; the general protocols for system preparation, convergence assessment, and validation transfer broadly across methods and systems.

Contents

I	Theoretical Foundations	4
1	Introduction	4
1.1	Why Free Energies Matter	4
1.2	The Fundamental Challenge	4
1.3	Three Strategies for Computing Free Energy Differences	5
2	Conceptual Development	5
2.1	Theoretical Foundations (1935-1999)	5
2.2	From Theory to Practice (1960s-2000s)	6
2.3	Modern Era (2000s-Present)	6
3	Theoretical Foundations	7
3.1	Free Energy and Partition Functions	7
3.2	Free Energy Differences as Probability Ratios	7
3.3	Phase Space and Ensemble Averages	8
3.4	Thermodynamic Cycles	8
3.5	The Alchemical Pathway	10

4	Free Energy Perturbation (FEP)	11
4.1	The Central Idea	11
4.2	The Zwanzig Equation	11
4.3	The Overlap Problem and Its Solution	12
5	Thermodynamic Integration (TI)	13
5.1	Motivation: An Alternative Formulation	13
5.2	The TI Equation	13
5.3	TI vs FEP: Why Derivatives Average Better	14
5.4	Practical Considerations	14
6	Bennett Acceptance Ratio (BAR)	15
6.1	The Bidirectional Advantage	15
6.2	Why BAR is Optimal	15
6.3	From BAR to MBAR	16
7	Nonequilibrium Methods	16
7.1	A Different Paradigm	16
7.2	The Jarzynski Equality	16
7.3	The Crooks Fluctuation Theorem	17
7.4	Practical Implementation	17
7.5	When to Use Nonequilibrium Methods	18
II	Practical Implementation	19
8	Setting Up FEP Calculations	19
8.1	Choosing λ Windows	19
8.2	Simulation Protocol	19
8.3	Soft-Core Potentials	20
8.4	What GROMACS Computes	20
9	Analysis and Validation	20
9.1	The pmx analyze Step	20
9.2	Convergence Diagnostics	20
9.2.1	Method-Specific Checks	20
9.2.2	Universal Diagnostics	21
9.3	Error Estimation	21
9.4	Troubleshooting	21
10	GROMACS/pmx Tutorial	21
10.1	Prerequisites and Installation	21
10.2	Workflow Overview	22
10.3	System Preparation	22
10.4	Lambda Windows and Simulations	23
10.5	Analysis and Quality Control	24
11	Case Study: CDK8 Kinase Inhibitors	24
11.1	The System	24
11.2	Protocol	25
11.3	Results and Lessons Learned	25
11.4	Practical Takeaways	25

A	Derivation of the Zwanzig Equation	29
B	Derivation of Thermodynamic Integration	29
C	Derivation of the Jarzynski Equality	30
D	Bennett Acceptance Ratio Derivation	31
E	Software Comparison Table	31
F	Glossary	31

Part I

Theoretical Foundations

After reading this part, you will understand the statistical mechanical principles that enable free energy calculations from molecular simulations, the mathematical basis of FEP, TI, and nonequilibrium methods, and why these approaches are theoretically equivalent yet differ in practical implementation.

1 Introduction

1.1 Why Free Energies Matter

In biomolecular systems, free energy is the fundamental quantity that determines equilibrium behavior. Whether a protein folds, whether a ligand binds, whether a mutation stabilizes or destabilizes a structure—all of these questions are answered by free energy differences. Consider drug discovery: we want to know which of two candidate ligands binds more tightly to a target protein. The strength of binding is quantified by the dissociation constant K_d , which is related to the binding free energy ΔG_{bind} :

$$K_d/C^0 = e^{\beta\Delta G_{bind}} \quad (1)$$

where $\beta = 1/kT$ and $C^0 = 1M$ is the standard state concentration. Stronger binding corresponds to smaller K_d and more negative ΔG_{bind} . A difference of just 1.4 kcal/mol (or $2.35 kT$ at 300 K) corresponds to a 10-fold change in K_d at room temperature. This is the difference between a promising drug candidate and an inactive compound.

Yet measuring free energies experimentally is challenging and time-consuming. If we could compute them accurately from first principles, we could:

- Screen thousands of compounds computationally before synthesis
- Understand why certain mutations affect protein stability
- Design better drugs by predicting which modifications improve binding
- Explain experimental observations at the molecular level

Computational free energy methods aim to deliver this capability by connecting atomic-level simulations to macroscopic thermodynamic observables.

1.2 The Fundamental Challenge

Free energy is defined in terms of the partition function:

$$F = -kT \ln Z, \quad Z = \int e^{-\beta H(\mathbf{x})} d\mathbf{x} \quad (2)$$

where $H(\mathbf{x})$ is the Hamiltonian (total energy) and the integral is over all possible configurations $\mathbf{x}$ of the system.

The problem: even for a modest protein-ligand system, the configuration space has tens of thousands of dimensions. Computing this integral directly is completely impossible. We need methods that can extract free energies from molecular simulations without exhaustively sampling this vast space.

1.3 Three Strategies for Computing Free Energy Differences

In practice, we want to compare two states: does ligand A or ligand B bind more tightly? Is the wild-type or mutant protein more stable? The challenge is that simulations of state A sample configurations relevant to A but rarely explore regions relevant to state B. This overlap problem—states occupying different regions of configuration space—is what alchemical methods address. The methods covered in this document represent three different strategies for tackling this challenge:

Free Energy Perturbation (FEP): Connect states A and B through an alchemical pathway with intermediate states. Compute the free energy difference using exponential averaging (Zwanzig equation) or optimal bidirectional estimators (BAR).

Thermodynamic Integration (TI): Use the same alchemical pathway, but compute derivatives $\partial F/\partial\lambda$ and integrate rather than using exponential averaging.

Nonequilibrium methods: Perform many independent trajectories that switch the system from A to B, measure the work along each path, and use the Jarzynski equality or Crooks theorem to extract the equilibrium free energy difference.

Each approach has strengths and weaknesses.

Part I develops the statistical mechanical framework underlying free energy calculations, covering FEP, TI, and nonequilibrium methods from first principles. We derive the key equations (Zwanzig, BAR, Jarzynski equality) and explain their connections and practical implications.

Part II addresses the simulation protocols, convergence assessment, and validation procedures needed for reliable calculations, including a complete GROMACS/pmx tutorial and lessons from a 54-edge CDK8 kinase benchmark study. Throughout, we emphasize the connection between theoretical rigor and practical computational workflows.

2 Conceptual Development

The methods covered in this document emerged through a series of theoretical breakthroughs, each addressing limitations of previous approaches. Understanding this progression clarifies why we have multiple methods and when each is most appropriate.

2.1 Theoretical Foundations (1935-1999)

Kirkwood (1935): John Kirkwood derived what we now call thermodynamic integration in his paper "Statistical Mechanics of Fluid Mixtures". He showed that free energy differences could be computed by integrating $\langle\partial H/\partial\lambda\rangle$ along a coupling parameter λ . This was purely theoretical work—digital computers didn't exist yet. Kirkwood's motivation was understanding liquid theories, not computing actual free energies for complex molecular systems.

Zwanzig (1954): Nearly two decades later, Robert Zwanzig published "High-Temperature Equation of State by a Perturbation Method." His elegant equation:

$$\Delta G = -kT \ln \langle e^{-\beta(H_B - H_A)} \rangle_A \quad (3)$$

represented a conceptual breakthrough: you could compute $G_B - G_A$ from a simulation of state A alone. However, Zwanzig himself recognized this was mainly useful for small perturbations—for large changes, the exponential average would be dominated by rare configurations that simulations couldn't adequately sample.

Bennett (1976): Charles Bennett derived the optimal estimator for combining forward and reverse free energy data in his 1976 paper. The method (now known as the Bennett Acceptance Ratio or BAR) was largely ignored for nearly two decades. At the time, few researchers were actively doing free energy calculations, and computational limitations made even basic FEP impractical, let alone optimized variants. The method's importance would only be recognized much later.

Jarzynski (1997) and Crooks (1999): Christopher Jarzynski published "Nonequilibrium Equality for Free Energy Differences" with a startling result:

$$e^{-\beta\Delta G} = \langle e^{-\beta W} \rangle \quad (4)$$

You could extract *equilibrium* free energies from *nonequilibrium* work measurements. Gavin Crooks extended the framework with his bidirectional fluctuation theorem, making nonequilibrium methods more practical. However, they haven't replaced equilibrium FEP for routine applications—the exponential averaging problem persists, and computational cost ends up similar.

2.2 From Theory to Practice (1960s-2000s)

For decades following Zwanzig's work, FEP and TI remained largely theoretical curiosities. Early computers were too slow for adequate sampling, and pioneering calculations in the 1960s-70s used tiny systems with severe convergence problems.

The 1980s brought improved workstations and the first biomolecular FEP attempts by Tembe & McCammon, Jorgensen, and others. Results were often unreliable, but the field learned about the overlap problem, equilibration requirements, and sampling challenges.

Several developments transformed FEP from fragile to robust:

Rediscovery of BAR (1990s): Researchers recognized Bennett's 1976 work provided a dramatically better estimator than one-directional FEP, significantly reducing sampling requirements for convergence.

Soft-core potentials (Beutler et al., 1994): When atoms appear or disappear, naive interpolation creates singularities. Soft-core potentials smooth these transitions, making previously impossible transformations routine.

Force field maturation: As biomolecular force fields (AMBER, CHARMM, GROMOS) improved through the 1990s-2000s, FEP results became quantitatively meaningful rather than qualitatively suggestive. Free energy calculations can only be as accurate as the underlying energy model.

MD algorithm improvements: Better integrators, pressure coupling, and long-range electrostatics made simulations faster and more reliable.

These advances were synergistic. Better algorithms reduced sampling requirements; successful predictions—FEP correlating with experimental binding affinities—attracted more researchers and drove further development.

2.3 Modern Era (2000s-Present)

The hardware revolution fundamentally changed molecular dynamics economics. Modern CPUs and especially GPUs deliver orders of magnitude more performance than previous generations. Calculations that once required expensive clusters now run on desktop workstations. A consumer GPU can complete multi-microsecond FEP campaigns for typical drug-like systems in days to weeks, enabling systematic validation at scales previously impractical.

As methods and hardware matured, pharmaceutical companies began adopting FEP in drug discovery pipelines. Both commercial platforms and academic tools demonstrated impressive accuracy in validation studies and blind challenges. Free energy calculations transitioned from specialized research techniques to routine computational tools.

Today, the theoretical foundations are thoroughly understood, algorithms are mature and well-validated, and computational power is accessible. The bottleneck has shifted from "Can we compute this?" to questions of convergence validation, optimal perturbation network design, and integration into broader drug discovery workflows. Current frontier work explores integration with machine learning—for force field development, enhanced sampling, and learned alchemical pathways.

Key Insight

Understanding this conceptual development helps appreciate why modern FEP works as it does. Zwanzig’s equation is theoretically exact but computationally challenging (exponential averaging). Bennett showed how to optimally combine data (BAR). Soft-core potentials made diverse transformations practical. Modern hardware made extensive sampling affordable. Each advance built on previous work, and success required progress on multiple fronts simultaneously. This is a mature field that continues evolving as new computational and theoretical tools become available.

3 Theoretical Foundations

Before diving into specific methods, we establish the common theoretical framework underlying all free energy calculations.

3.1 Free Energy and Partition Functions

In statistical mechanics, all thermodynamic quantities derive from the partition function. For a system in thermal equilibrium at temperature T :

$$Z = \sum_i e^{-\beta E_i} \quad (\text{discrete states}) \quad (5)$$

or

$$Z = \int e^{-\beta H(\mathbf{x})} d\mathbf{x} \quad (\text{continuous phase space}) \quad (6)$$

where $\beta = 1/kT$, E_i or $H(\mathbf{x})$ is the energy of state i or configuration $\mathbf{x}$, and the sum/integral runs over all accessible states/configurations.

The Helmholtz free energy (constant volume, temperature) is:

$$F = -kT \ln Z \quad (7)$$

For biomolecular systems, we typically work at constant pressure and temperature, using the Gibbs free energy:

$$G = F + PV \approx F \quad (\text{pressure term usually negligible}) \quad (8)$$

In this document, we’ll use F and G interchangeably, as is common in the literature.

3.2 Free Energy Differences as Probability Ratios

Before discussing alchemical transformations, we need to understand what free energy differences mean at a fundamental level. Consider a system with **fixed molecular composition** that can exist in two distinct macroscopic states—for example, a protein in folded vs unfolded conformations, or different binding modes of the same ligand.

The free energy difference between these states is:

$$\Delta G = G_1 - G_0 = -kT \ln \frac{Z_1}{Z_0} \quad (9)$$

This seemingly abstract ratio of partition functions has a direct physical interpretation. In the canonical ensemble, the probability that the system occupies a particular microstate i with energy E_i is:

$$p_i = \frac{e^{-\beta E_i}}{Z} \quad (10)$$

where $Z = \sum_i e^{-\beta E_i}$ is the total partition function summed over all possible configurations.

Consider macrostate 0, which comprises many microstates. The total probability of finding the system in macrostate 0 is:

$$P_0 = \sum_{i \in \text{state 0}} p_i = \frac{1}{Z} \sum_{i \in \text{state 0}} e^{-\beta E_i} = \frac{Z_0}{Z} \quad (11)$$

where Z_0 is the restricted partition function for state 0.

Similarly, $P_1 = Z_1/Z$. Therefore:

$$\frac{P_1}{P_0} = \frac{Z_1}{Z_0} = e^{-\beta \Delta G} \quad (12)$$

****Free energy differences determine probability ratios.**** A system at equilibrium spends more time in the lower free energy state.

Key Insight

This connection between free energy and probability is fundamental. A large ΔG means one state is exponentially more probable than the other. If state 1 is much more favorable than state 0 (e.g., $P_1/P_0 = 10^{10}$), then simulations starting in state 1 will essentially never spontaneously visit state 0. This is why naive approaches fail for large free energy differences—we cannot directly sample exponentially rare transitions.

3.3 Phase Space and Ensemble Averages

For a system with N atoms, the phase space is $6N$ -dimensional: 3 position coordinates and 3 momentum coordinates for each atom. A microstate $\mathbf{x}$ specifies all positions and momenta.

An observable $A(\mathbf{x})$ (like energy, or distance between atoms) has an ensemble average:

$$\langle A \rangle = \frac{\int A(\mathbf{x}) e^{-\beta H(\mathbf{x})} d\mathbf{x}}{\int e^{-\beta H(\mathbf{x})} d\mathbf{x}} = \frac{1}{Z} \int A(\mathbf{x}) e^{-\beta H(\mathbf{x})} d\mathbf{x} \quad (13)$$

In molecular dynamics simulations, we approximate this by time averaging:

$$\langle A \rangle \approx \frac{1}{M} \sum_{i=1}^M A(\mathbf{x}_i) \quad (14)$$

where $\mathbf{x}_i$ are configurations sampled from the trajectory.

This is justified by the **ergodic hypothesis**: given enough time, a system will explore all accessible regions of phase space with frequency proportional to their Boltzmann weight. In practice, ergodicity is often violated (systems get trapped in local minima), which is a major source of convergence problems.

3.4 Thermodynamic Cycles

Since free energy is a state function, ΔG is path-independent—we can compute it along any convenient route connecting two states. We illustrate this using protein-ligand binding, though the same principles apply to protein stability, solvation, or membrane permeability calculations.

Consider a ligand binding to a protein in solution:



The binding free energy ΔG_{bind} quantifies this process's favorability; a negative value means spontaneous binding, and its magnitude determines the dissociation constant: $\Delta G_{bind} = RT \ln K_d$.

For drug discovery, we compare two ligands A and B. Their absolute binding free energies are:

$$\Delta G_{bind,A} = G(\text{complex}_A) - G(\text{protein}) - G(\text{ligand A}) \quad (16)$$


$$\Delta G_{bind,B} = G(\text{complex}_B) - G(\text{protein}) - G(\text{ligand B}) \quad (17)$$

where all species are in solution. The relative binding free energy is:

$$\Delta\Delta G_{bind} = \Delta G_{bind,B} - \Delta G_{bind,A} \quad (18)$$

Rather than computing absolute binding free energies directly (which requires sampling rare spontaneous binding events), we use a thermodynamic cycle with two alchemical transformations:

- ΔG_{bound} : transforming A→B while bound to protein
- ΔG_{free} : transforming A→B in solution



Going around the cycle gives:

$$\Delta G_{bind,A} + \Delta G_{bound} - \Delta G_{bind,B} - \Delta G_{free} = 0 \quad (19)$$

Rearranging:

$$\Delta\Delta G_{bind} = \Delta G_{bind,B} - \Delta G_{bind,A} = \Delta G_{bound} - \Delta G_{free} \quad (20)$$

****Key result:**** We compute relative binding affinity by alchemically transforming A→B in the binding site and in solution separately, then taking the difference.

Key Insight

Absolute binding free energies require sampling both bound and unbound states. The unbound ligand can be anywhere in solution, making the bound state exponentially rare. Relative free energies sidestep this: we transform ligand A to B in two fixed locations (binding site and bulk solvent), never sampling the binding/unbinding event itself. The alchemical transformation A→B is computationally tractable because we maintain structural similarity throughout the pathway. In drug discovery, relative binding affinities are usually what we want anyway.




Figure 1: Alchemical transformation pathway connecting states A and B through intermediate λ states. Each λ value defines a distinct thermodynamic state with its own partition function and free energy.

3.5 The Alchemical Pathway

The thermodynamic cycle tells us we need to compute ΔG_{bound} and ΔG_{free} —the free energy of transforming ligand A into ligand B in two different environments. But how do we actually compute the free energy difference between two chemically distinct molecules?

The key trick in FEP and TI is introducing an artificial coupling parameter $\lambda \in [0, 1]$ that continuously connects state A ($\lambda = 0$) to state B ($\lambda = 1$). We construct a λ -dependent Hamiltonian:

$$H(\mathbf{x}, \lambda) \quad (21)$$

with boundary conditions $H(\mathbf{x}, 0) = H_A(\mathbf{x})$ and $H(\mathbf{x}, 1) = H_B(\mathbf{x})$.

The simplest choice is linear interpolation:

$$H(\mathbf{x}, \lambda) = (1 - \lambda)H_A(\mathbf{x}) + \lambda H_B(\mathbf{x}) \quad (22)$$

However, this naive form creates problems when atoms appear or disappear (infinite forces at vanishing distances). In practice, more sophisticated schemes are used: separate λ parameters for electrostatics and van der Waals interactions, and soft-core potentials that smooth out singularities. We cover these details in Part II.

By introducing this continuous pathway, we can compute the free energy difference $\Delta G = G_B - G_A$ using statistical mechanics. The methods covered in this document—FEP, TI, and nonequilibrium approaches—represent different strategies for extracting ΔG from simulations along this λ pathway.

Key Insight

Why "alchemical"? Because intermediate λ values correspond to non-physical states. If A is a hydrogen atom and B is a chlorine atom, then $\lambda = 0.5$ represents something that's half H and half Cl—not an actual chemical species.

This is perfectly acceptable. We're using these intermediate states purely as a mathematical tool. Thermodynamics guarantees that ΔG depends only on the endpoints, not the path taken between them. We can use any convenient pathway, even an unphysical one.

The art of free energy calculations lies in designing λ pathways that are numerically well-behaved: good overlap between adjacent windows, smooth derivatives $\partial H/\partial \lambda$, and no singularities. These pathways may be chemically nonsensical, but they must be statistically tractable.

4 Free Energy Perturbation (FEP)

The term "FEP" has two common meanings: narrowly, the Zwanzig exponential formula specifically; broadly, any alchemical free energy calculation using λ -coupling. We use "FEP" to refer to the Zwanzig equation. Modern production calculations typically use more sophisticated estimators like BAR or MBAR, covered in later sections.

4.1 The Central Idea

Recall from the previous section that we want to compute free energy differences between chemically distinct molecules—for example, ΔG_{bound} for transforming ligand A into ligand B in the protein binding site. The alchemical pathway provides a continuous transformation via the coupling parameter λ , but how do we actually extract ΔG from simulations?

FEP's fundamental insight: ***sample configurations from one state and evaluate their energies in another state***. This perturbation approach avoids computing absolute free energies entirely.

4.2 The Zwanzig Equation

Theorem 4.1 (Zwanzig, 1954). *The free energy difference between two states can be computed from a simulation of state A alone:*

$$\Delta G = G_B - G_A = -kT \ln \left\langle e^{-\beta(H_B - H_A)} \right\rangle_A \quad (23)$$

where $\beta = 1/kT$, H_A and H_B are the Hamiltonians of the two states, and $\langle \dots \rangle_A$ denotes an ensemble average over configurations sampled from equilibrium at state A.

See Appendix A for the derivation.

Key Insight

The Zwanzig equation says: run a simulation in state A, and for each configuration sampled, compute the energy difference $\Delta H = H_B - H_A$. Then compute the exponential average $\langle e^{-\beta \Delta H} \rangle_A$.

Notice something remarkable: we compute ΔG from sampling only state A. We never simulate state B explicitly—we just evaluate H_B at configurations drawn from A's distribution.

==

4.3 The Overlap Problem and Its Solution

The Zwanzig equation is mathematically exact, but computationally treacherous. The exponential makes it extremely sensitive to rare events.

Consider a concrete example: you sample 1000 configurations from state A and compute $\Delta H = H_B - H_A$ for each. Suppose most configurations have $\Delta H \approx 5$ kcal/mol, but a few rare configurations have $\Delta H \approx -2$ kcal/mol.

At room temperature ($\beta \approx 1.68$ (kcal/mol) $^{-1}$):

- Typical configurations: $e^{-\beta \cdot 5} \approx e^{-8.4} \approx 2 \times 10^{-4}$
- Rare configurations: $e^{-\beta \cdot (-2)} \approx e^{3.4} \approx 30$

The rare configurations contribute about 10^5 times more per configuration than typical ones. Even if they represent only a few percent of your samples, they dominate $\langle e^{-\beta \Delta H} \rangle_A$. If you don't sample these rare low- ΔH configurations adequately, your free energy estimate will be severely biased.

Key Insight

Exponential averaging amplifies contributions from rare low- ΔH configurations. Even moderate energy differences lead to factors of 10^5 or more in relative weights. Unless you adequately sample these rare events, your free energy estimate will be biased.

This is the ****overlap problem****: you need configurations where states A and B have similar energies—where their energy distributions overlap. If A and B differ substantially, this overlap region is exponentially rare in A's equilibrium distribution.

Common Pitfall

A common misconception: "I ran the simulation for a long time, so it should converge." Wrong! If states A and B are too different, even extensive sampling of A won't help. The relevant configurations—those where $H_B \approx H_A$ —are exponentially suppressed in state A's Boltzmann distribution. More sampling of the wrong region doesn't solve the problem.

Example: Computing absolute binding free energy by simulating the unbound ligand and evaluating bound-state energies. Configurations where the ligand happens to be positioned correctly in the binding site are astronomically rare. No amount of unbound-state sampling will adequately capture them.

The solution: gradual transformations. Don't transform A to B in one giant leap. Instead, introduce intermediate states using the alchemical pathway from Section 3.5.

Using the λ -parameterized Hamiltonian:

$$H(\mathbf{x}, \lambda) = (1 - \lambda)H_A(\mathbf{x}) + \lambda H_B(\mathbf{x}) \quad (24)$$

we divide the transformation into N windows: $0 = \lambda_0 < \lambda_1 < \dots < \lambda_N = 1$.

The total free energy change is:

$$\Delta G = F(\lambda_N) - F(\lambda_0) = \sum_{i=0}^{N-1} [F(\lambda_{i+1}) - F(\lambda_i)] \quad (25)$$

Each window contributes:

$$\Delta G_i = F(\lambda_{i+1}) - F(\lambda_i) = -kT \ln \left\langle e^{-\beta[H(\lambda_{i+1}) - H(\lambda_i)]} \right\rangle_{\lambda_i} \quad (26)$$

If the λ spacing is sufficiently fine, adjacent windows have good overlap. Each individual Zwanzig step becomes well-behaved, even though the overall A→B transformation would be intractable in one step.

Key Insight

How many λ windows do you need? It depends on how different A and B are. A rule of thumb: aim for $\Delta G_i \lesssim 2\text{--}3$ kT per window. Larger steps risk poor overlap; smaller steps waste computational effort.

For typical drug-like ligand transformations, 10–20 windows often suffice. Transformations involving charge changes or atom deletions may require 30–50 windows. We discuss practical guidelines in Part II.

5 Thermodynamic Integration (TI)

5.1 Motivation: An Alternative Formulation

We’ve seen that FEP computes free energy differences via exponential averaging:

$$\Delta G = -kT \ln \langle e^{-\beta \Delta H} \rangle \quad (27)$$

This works when adjacent λ windows have good overlap, but the exponential remains problematic—it amplifies contributions from rare configurations. Even with optimal estimators like BAR, we’re managing the exponential’s statistical behavior.

Is there a fundamentally different approach using the same alchemical pathway? Yes: instead of averaging exponentials of energy differences, we can integrate averages of energy derivatives.

5.2 The TI Equation

Theorem 5.1 (Thermodynamic Integration, Kirkwood 1935). *The free energy difference along an alchemical pathway can be computed by integrating the ensemble average of the energy derivative:*

$$\Delta G = G(1) - G(0) = \int_0^1 \left\langle \frac{\partial H(\lambda)}{\partial \lambda} \right\rangle_{\lambda} d\lambda \quad (28)$$

where $\langle \cdots \rangle_{\lambda}$ denotes averaging over configurations sampled at intermediate state λ .

See Appendix B for the derivation.

The equation is remarkably intuitive: it says that the rate at which free energy changes, $\frac{dG}{d\lambda}$, equals the ensemble average of how the Hamiltonian changes, $\left\langle \frac{\partial H(\lambda)}{\partial \lambda} \right\rangle_{\lambda}$. Integrate this rate from $\lambda = 0$ to $\lambda = 1$, and you recover the total free energy difference—it’s the fundamental theorem of calculus applied to thermodynamics.

Physically, $\left\langle \frac{\partial H}{\partial \lambda} \right\rangle_{\lambda}$ represents the “thermodynamic force” driving the transformation from A to B. A gentle slope indicates an easy transformation; a steep slope reveals where the free energy changes rapidly. This makes the $\partial H/\partial \lambda$ vs λ plot invaluable diagnostically—you can literally see which regions need denser λ sampling.

In practice: simulate at several intermediate λ values, compute the ensemble average of $\partial H/\partial \lambda$ at each window, then numerically integrate over λ .

5.3 TI vs FEP: Why Derivatives Average Better

TI and FEP compute the same quantity using different formulations:

$$\text{FEP: } \Delta G = -kT \ln \left\langle e^{-\beta(H_B - H_A)} \right\rangle_A \quad (29)$$

$$\text{TI: } \Delta G = \int_0^1 \left\langle \frac{\partial H(\lambda)}{\partial \lambda} \right\rangle_\lambda d\lambda \quad (30)$$

Both are mathematically exact with perfect sampling, but TI has a key statistical advantage: it uses linear averaging of derivatives instead of exponential averaging of energy differences.

Consider energy differences between adjacent λ windows where most values cluster around 5 kcal/mol, with rare outliers at -2 kcal/mol.

FEP exponential average:

- Typical: $e^{-\beta \cdot 5} \approx 2 \times 10^{-4}$
- Rare: $e^{-\beta \cdot (-2)} \approx 30$
- Factor difference: $\sim 10^5$ (rare events dominate!)

TI linear average of $\partial H/\partial \lambda$:

- Rare events contribute proportionally to their occurrence
- No exponential amplification
- Variance affects uncertainty, not the estimator itself

Mathematically, the cumulant expansion reveals why:

$$\ln \langle e^X \rangle \approx \langle X \rangle + \frac{1}{2} \text{Var}(X) + \dots \quad (31)$$

The FEP estimator $-kT \ln \langle e^{-\beta \Delta H} \rangle$ depends on both mean *and* variance of ΔH . Large variance (poor overlap) introduces systematic bias. The TI estimator $\langle \partial H/\partial \lambda \rangle$ depends only on the mean—variance affects error bars, but the estimator remains unbiased.

5.4 Practical Considerations

When is TI advantageous?

- Overlap between adjacent λ windows is marginal
- You want diagnostic information on where free energy changes rapidly
- You're implementing from scratch and prefer simpler analysis

Key practical point: Both methods require identical simulation setups—the same λ windows and sampling. Modern MD packages (GROMACS, AMBER) output both $\partial H/\partial \lambda$ (for TI) and energy differences to neighboring windows (for FEP/BAR) simultaneously. You can compute both estimates from a single set of simulations.

The difference is purely in post-processing: TI integrates derivatives; BAR uses exponential reweighting optimally. With proper λ spacing, both give essentially identical results. The choice often comes down to which diagnostic plots you find more interpretable.

Key Insight

The "TI vs FEP" debate is largely obsolete. You run simulations at multiple λ windows and compute both estimates. If they agree, you're converged. If they disagree substantially, you have a sampling problem—neither estimator is reliable yet.

The practical questions that matter:

- How many λ windows?
- How much sampling per window?
- How do I diagnose convergence?

These apply equally to TI and FEP. We address them in Part II.

6 Bennett Acceptance Ratio (BAR)

FEP and TI are theoretical frameworks for computing free energies from molecular simulations. BAR is an optimal statistical method for *estimating* free energies from finite simulation data.

6.1 The Bidirectional Advantage

The Zwanzig equation uses only forward simulations (sampling state A, evaluating energies in state B). But if we simulate both states A and B, can we extract more information?

Yes. Charles Bennett showed in 1976 that there's an optimal way to combine forward and reverse data, achieving lower variance than either direction alone.

Theorem 6.1 (Bennett Acceptance Ratio). *Given n_A samples from state A and n_B samples from state B, the optimal free energy estimate satisfies:*

$$\sum_{j=1}^{n_A} f(\beta(H_B(\mathbf{x}_j^A) - H_A(\mathbf{x}_j^A)) + C) = \sum_{k=1}^{n_B} f(\beta(H_A(\mathbf{x}_k^B) - H_B(\mathbf{x}_k^B)) - C) \quad (32)$$

where $f(x) = 1/(1 + e^x)$ is the Fermi function, and $C = \beta\Delta G + \ln(n_B/n_A)$.

This equation is solved iteratively for ΔG .

See Appendix ?? for the derivation.

6.2 Why BAR is Optimal

Key Insight

BAR combines forward ($A \rightarrow B$) and reverse ($B \rightarrow A$) information optimally. The Fermi function weights configurations based on their contribution to the overlap region where both states are relevant.

Mathematically, BAR is the maximum likelihood estimator of ΔG given samples from both states. No other estimator achieves lower variance with the same data.

Practically: if forward Zwanzig gives one estimate and reverse gives another (indicating poor overlap), BAR finds the statistically optimal compromise. This is why BAR converges much faster than one-directional FEP.

6.3 From BAR to MBAR

In typical FEP calculations with many λ windows (20-50 intermediate states), we use the Multistate Bennett Acceptance Ratio (MBAR). MBAR generalizes BAR to simultaneously optimize across all states, extracting maximum information from the entire dataset.

Modern software implements MBAR by default:

- GROMACS with `gmx bar` or `pmx analyze`
- AMBER with `alchemical-analysis`
- Python tools like `pymbar`

These tools also provide diagnostics—overlap matrices, convergence plots, error estimates—to verify calculation quality. We cover practical usage in Part II.

Bottom line: Always use BAR/MBAR when you have data from multiple λ windows. It’s the statistically optimal estimator and is now standard in all modern FEP workflows.

7 Nonequilibrium Methods

7.1 A Different Paradigm

FEP, TI, and BAR all require equilibrium sampling at fixed λ values. You simulate at each window until the system equilibrates, then collect statistics from the equilibrium distribution.

Nonequilibrium methods take a fundamentally different approach: ****perform many independent switching trajectories that rapidly drive the system from A to B, measure the work along each path, and extract the equilibrium free energy from the work distribution****.

This seems paradoxical—doesn’t irreversible work always exceed the equilibrium free energy? Yes, for any single trajectory. But by averaging over many trajectories in a specific way, exact equalities discovered by Jarzynski (1997) and Crooks (1999) allow recovery of the equilibrium ΔG .

7.2 The Jarzynski Equality

Consider a system with time-dependent Hamiltonian $H(\mathbf{x}, t)$ where the coupling parameter $\lambda(t)$ is driven from 0 to 1 over time interval $[0, \tau]$. The system starts in equilibrium at $\lambda = 0$ (state A), and we switch to $\lambda = 1$ (state B) according to some protocol $\lambda(t)$.

The work performed along a single trajectory is:

$$W = \int_0^\tau \frac{\partial H(\mathbf{x}(t), t)}{\partial t} dt = \int_0^\tau \frac{\partial H}{\partial \lambda} \frac{d\lambda}{dt} dt \quad (33)$$

Different trajectories—different initial configurations and different stochastic fluctuations—yield different work values. The work distribution $P(W)$ encodes the free energy difference.

Theorem 7.1 (Jarzynski Equality, 1997). *For a system driven from equilibrium state A by a time-dependent Hamiltonian:*

$$\Delta G = -kT \ln \langle e^{-\beta W} \rangle \quad (34)$$

where the average is over all nonequilibrium switching trajectories starting from equilibrium at A.

See Appendix C for the derivation.

Key Insight

Individual trajectories are irreversible and typically perform excess work: $W > \Delta G$. Yet the exponentially-weighted average $\langle e^{-\beta W} \rangle$ equals $e^{-\beta \Delta G}$ exactly.

Why? Rare trajectories with $W < \Delta G$ (where thermal fluctuations help the transformation) are exponentially upweighted, exactly compensating for typical high-work trajectories. This is not a second law violation—the second law says $\langle W \rangle \geq \Delta G$. Jarzynski gives you ΔG , not $\langle W \rangle$.

Common Pitfall

”Can I do one fast switching trajectory and get ΔG from that single W value?”

No! Individual trajectories typically have $W \gg \Delta G$ due to dissipated work. You need many trajectories (hundreds to thousands) to sample the work distribution. The exponential average is dominated by rare low- W trajectories—the same exponential averaging problem as in FEP, just manifested differently.

7.3 The Crooks Fluctuation Theorem

Jarzynski uses only forward ($A \rightarrow B$) trajectories. Crooks (1999) discovered a symmetry relation when reverse ($B \rightarrow A$) trajectories are included.

Theorem 7.2 (Crooks Fluctuation Theorem, 1999). *Let $P_F(W)$ be the work distribution for forward trajectories and $P_R(W)$ for reverse trajectories. Then:*

$$\frac{P_F(W)}{P_R(-W)} = e^{\beta(W - \Delta G)} \quad (35)$$

At the crossing point where $P_F(W^*) = P_R(-W^*)$, we have $W^* = \Delta G$ directly—no exponential averaging required.

Key Insight

The Crooks theorem provides:

- Visual validation: plot $P_F(W)$ and $P_R(-W)$ to verify they satisfy the predicted relationship
- Direct ΔG estimation from the distribution crossing point
- Convergence diagnostics: deviations from the Crooks relation indicate inadequate sampling

Bidirectional protocols (combining forward and reverse via Crooks or BAR-like estimators) converge faster than unidirectional Jarzynski.

7.4 Practical Implementation

Switching protocol $\lambda(t)$: Linear switching ($\lambda(t) = t/\tau$) is simplest but not optimal. Smoother protocols (slow acceleration at endpoints) can reduce dissipated work, though this is problem-dependent.

Switching time τ : Fast switching (nanoseconds) produces far-from-equilibrium trajectories with large work fluctuations—many trajectories needed. Slow switching (microseconds) gives smaller fluctuations but longer individual runs. Total cost (switching time $\times$ number of trajectories) ends up comparable to equilibrium FEP.

Typical workflow:

1. Run equilibrium simulation at state A, save configurations
2. Launch switching trajectories $A \rightarrow B$ from each configuration, record work W_F
3. Run equilibrium simulation at state B, save configurations
4. Launch switching trajectories $B \rightarrow A$ from each configuration, record work W_R
5. Combine W_F and W_R via Crooks relation or BAR-like estimator

Modern tools like pmx with nonequilibrium support automate this workflow.

7.5 When to Use Nonequilibrium Methods

Given comparable computational cost to equilibrium FEP, when choose nonequilibrium?

Advantages:

- Systems with high barriers or slow conformational changes where equilibrium sampling gets trapped
- Fast switching can traverse barriers that equilibrium simulations struggle with
- Direct connection to single-molecule experiments (optical tweezers, AFM)
- Compatibility with enhanced sampling techniques

Current limitations:

- Exponential averaging problem persists (rare low-work trajectories dominate)
- Sensitivity to switching protocol and speed
- Less mature software and less standardized workflows than equilibrium FEP
- More difficult convergence assessment

For routine applications like drug discovery (relative binding free energies for compound ranking), equilibrium FEP/BAR remains the standard. Nonequilibrium methods excel for challenging systems where equilibrium sampling fails.

Key Insight

Nonequilibrium methods are powerful for specific problems—especially those with kinetic barriers or slow relaxation. They’re actively used in research and continue to develop. However, for production drug discovery workflows, equilibrium FEP/BAR is the workhorse. Nonequilibrium methods are “advanced specialty tools” rather than “everyday standard practice.”

This may shift as methodology and software mature. The theory is sound—it’s a question of implementation, best practices, and user-friendliness.

Part II

Practical Implementation

8 Setting Up FEP Calculations

8.1 Choosing λ Windows

Number of windows:

- Minimum: 10-15 windows for simple transformations
- Standard: 20-25 windows for most applications
- Dense sampling: 30-40 windows for charge-changing or difficult transformations

Window spacing:

- Denser near endpoints ($\lambda = 0$ and $\lambda = 1$) where overlap is typically poorest
- Example spacing for 21 windows:

= 0.00, 0.05, 0.10, 0.20, 0.30, 0.40, 0.50,
0.60, 0.70, 0.80, 0.90, 0.95, 1.00

Adaptive strategies:

1. Run initial calculation with coarse spacing (10-12 windows)
2. Examine overlap matrix and $\langle \partial H / \partial \lambda \rangle$ curve
3. Add windows in regions of poor overlap or high curvature
4. Re-run only the new windows (old data still valid)

8.2 Simulation Protocol

Equilibration:

- NVT equilibration: 100-500 ps (temperature coupling)
- NPT equilibration: 500 ps - 2 ns (pressure equilibration)
- Check: energies stable, box dimensions equilibrated, RMSD plateaued

Production:

- Standard: 5-20 ns per window
- Challenging systems: 50-100 ns per window
- Goal: sufficient sampling for converged ensemble averages

Time step:

- 2 fs with bond constraints (LINCS/SHAKE)
- 1 fs if experiencing instabilities

8.3 Soft-Core Potentials

One technical detail: when transforming between chemically different molecules (e.g., changing an atom from carbon to hydrogen), direct linear interpolation creates singularities.

The problem: At $\lambda \approx 0$, the nearly-disappeared atom from state B can overlap with other atoms, causing energy to blow up.

The solution: Soft-core potentials that smoothly "turn off" atoms without creating singularities.

GROMACS soft-core parameters:

```
sc-alpha      = 0.5      ; Soft-core alpha parameter
sc-power      = 1        ; Soft-core power
sc-sigma      = 0.3      ; Soft-core sigma
```

These are implemented automatically when you specify topology changes. The math is technical, but the concept is: smooth interpolation without unphysical overlaps.

8.4 What GROMACS Computes

During a simulation at λ_i :

- GROMACS samples configurations from the equilibrium distribution at $H(\lambda_i)$
- For each configuration, it evaluates $H(\lambda)$ at neighboring λ values
- It outputs $dH/d\lambda$ (used for TI) and energy differences to neighboring states (used for FEP/BAR)

These outputs (in `dhd1.xvg` files) contain all information needed to apply TI, FEP, or BAR.

9 Analysis and Validation

9.1 The pmx analyze Step

After running all λ windows, you use `pmx analyze` to:

1. Read the energy differences from all windows
2. Apply BAR between each pair of consecutive windows: $\Delta G_i = F(\lambda_{i+1}) - F(\lambda_i)$
3. Sum up: $\Delta G_{\text{total}} = \sum_i \Delta G_i$
4. Estimate errors (typically via bootstrapping)

The output is your free energy difference ΔG_{B-A} with error bars.

9.2 Convergence Diagnostics

9.2.1 Method-Specific Checks

FEP/BAR: Overlap matrices

For each pair of adjacent λ windows, BAR computes an "overlap" metric (related to the probability of accepting a configuration from one state in the other).

Rule of thumb: Overlap should be > 0.03 for reliable results. If overlap < 0.01 , you need denser λ spacing or more sampling.

TI: Integrand smoothness

Plot $\langle \partial H / \partial \lambda \rangle$ vs λ . It should be:

- Smooth (no wild oscillations)

- Consistent between forward and reverse directions
- Well-resolved (enough λ points to capture the shape)

If the curve is noisy or shows hysteresis, you need more sampling or better λ spacing.

9.2.2 Universal Diagnostics

Cycle closure:

The most powerful validation: construct thermodynamic cycles and check closure.

For a cycle $A \rightarrow B \rightarrow C \rightarrow A$:

$$\Delta G_{AB} + \Delta G_{BC} + \Delta G_{CA} \stackrel{?}{=} 0 \quad (36)$$

Deviations indicate unconverged sampling somewhere in the cycle.

Acceptable closure: < 0.5 kcal/mol for well-converged calculations. If closure error > 1 kcal/mol, investigate.

Hysteresis (forward vs reverse):

If you compute free energy in both directions ($A \rightarrow B$ and $B \rightarrow A$), they should agree:

$$\Delta G_{AB} \stackrel{?}{=} -\Delta G_{BA} \quad (37)$$

Significant hysteresis (> 0.5 kcal/mol) indicates:

- Insufficient equilibration at some λ windows
- Barriers that cause slow relaxation
- Need for longer sampling or enhanced sampling techniques

Time evolution:

Plot ΔG vs simulation time. It should:

- Converge to a stable value
- Have error bars that shrink with more data
- Not show systematic drift

9.3 Error Estimation

Bootstrapping: Resample trajectory data with replacement, recompute ΔG many times (1000-10000 iterations), examine the distribution. Standard deviation = error estimate. Implemented in pmx and alchemlyb.

Blocking analysis: For time-correlated MD data, divide trajectory into blocks larger than autocorrelation time, compute ΔG for each block, estimate error from block-to-block variation.

Multiple independent runs: Gold standard: run 3-5 independent calculations from different initial conditions. If all agree within error bars, high confidence in convergence. Expensive ($5\times$ cost) but definitive.

9.4 Troubleshooting

10 GROMACS/pmx Tutorial

10.1 Prerequisites and Installation

Required software:

Symptom	Possible Cause	Solution
Poor overlap (< 0.01)	Lambda spacing too coarse	Add intermediate windows
Large cycle closure errors	One edge unconverged	Increase sampling or windows
Hysteresis	Insufficient equilibration	Extend equilibration time
Noisy $dH/d\lambda$	Rare events, barriers	More sampling or restraints
TI and BAR disagree	Convergence problem	Check overlap, extend sampling

Table 1: Common problems and solutions

- GROMACS 2020+ (GPU-enabled recommended)
- Python 3.7+ with pmx
- Molecular viewer (PyMOL/VMD)

Quick installation:

```

1 # GROMACS via conda (easiest)
2 conda install -c conda-forge gromacs
3
4 # pmx
5 pip install pmx
6
7 # Verify
8 gmx -version
9 pmx --help

```

For GPU support, compile GROMACS from source with CUDA enabled. See gromacs.org for detailed instructions.

10.2 Workflow Overview

Computing $\Delta\Delta G$ for transforming ligand A to ligand B:

1. Prepare structures (protein + ligands)
2. Create hybrid topology with pmx
3. Build and solvate systems (complex and solvent)
4. Energy minimize
5. Set up λ windows
6. Run simulations (equilibration + production)
7. Analyze with pmx

10.3 System Preparation

Protein and ligand setup:

```

1 # Clean protein
2 grep -v "HOH\|HETATM" protein_raw.pdb > protein_clean.pdb
3
4 # Generate topology
5 gmx pdb2gmx -f protein_clean.pdb -o protein.gro \
6     -water tip3p -ff amber99sb-ildn
7

```

```

8 # Assume you have ligandA.pdb, ligandB.pdb with force field parameters
9 # (ligandA.itp, ligandB.itp from GAFF/CGenFF/ATB)

```

Create hybrid topology:

```

1 # Automatic atom mapping
2 pmx atomMapping -i1 ligandA.pdb -i2 ligandB.pdb -o mapping.txt
3
4 # Generate hybrid structure
5 pmx mutate -f ligandA.gro -fB ligandB.gro \
6     --itp ligandA.itp --mapping mapping.txt \
7     -o ligandAB.gro --itp_out ligandAB.itp

```

Build systems (complex and solvent):

```

1 # Complex: combine protein + hybrid ligand
2 cat protein.gro ligandAB.gro > complex.gro
3
4 # Solvate in box
5 gmx editconf -f complex.gro -o complex_box.gro -c -d 1.2 -bt cubic
6 gmx solvate -cp complex_box.gro -cs spc216.gro -o complex_solv.gro -p
   topol.top
7
8 # Add ions (neutralize + 0.15 M NaCl)
9 gmx grompp -f ions.mdp -c complex_solv.gro -p topol.top -o ions.tpr
10 gmx genion -s ions.tpr -o complex_ions.gro -p topol.top \
11     -pname NA -nname CL -neutral -conc 0.15
12
13 # Energy minimization
14 gmx grompp -f em.mdp -c complex_ions.gro -p topol.top -o em.tpr
15 gmx mdrun -deffnm em
16
17 # Repeat for solvent leg (ligand in water)

```

10.4 Lambda Windows and Simulations

Key FEP parameters in mdp file:

```

1 free-energy          = yes
2 init-lambda-state   = 0      ; Changes for each window
3
4 ; 21 lambda windows
5 fep-lambdas          = 0.00 0.05 0.10 0.15 0.20 0.25 0.30 0.35 0.40
6                      0.45 0.50 0.55 0.60 0.65 0.70 0.75 0.80 0.85
7                      0.90 0.95 1.00
8
9 ; Soft-core (avoid singularities)
10 sc-alpha             = 0.5
11 sc-power             = 1
12 sc-sigma             = 0.3
13
14 nstdhdl              = 100   ; Output frequency

```

Run equilibration and production:

```

1 for lambda in {0..20}; do
2     # Generate TPR
3     sed "s/init-lambda-state_000=0/init-lambda-state_000=${lambda}/" \
4         fep_template.mdp > fep_${lambda}.mdp

```

```

5      gmx grompp -f fep_{$lambda}.mdp -c em.gro -p topol.top -o fep_{$
      lambda}.tpr
6
7      # NVT + NPT equilibration (100 ps each)
8      gmx mdrun -deffnm nvt_{$lambda} -s fep_{$lambda}.tpr
9      gmx grompp -f npt.mdp -c nvt_{$lambda}.gro -t nvt_{$lambda}.cpt \
10     -p topol.top -o npt_{$lambda}.tpr
11     gmx mdrun -deffnm npt_{$lambda}
12
13     # Production (20 ns)
14     gmx grompp -f prod.mdp -c npt_{$lambda}.gro -t npt_{$lambda}.cpt \
15     -p topol.top -o prod_{$lambda}.tpr
16     gmx mdrun -deffnm prod_{$lambda}
17 done

```

Time estimate: 3-4 hours per window on RTX 4090, 80 hours total per edge.

10.5 Analysis and Quality Control

Analyze with pmx:

```

1 # Complex leg
2 pmx analyse -fA complex/prod_*/dhdl.xvg -o results_complex.txt
3
4 # Solvent leg
5 pmx analyse -fA solvent/prod_*/dhdl.xvg -o results_solvent.txt

```

Example output:

Complex: DG = -5.23 ± 0.31 kcal/mol
 Solvent: DG = -2.08 ± 0.18 kcal/mol
 Relative: DDG = -3.15 ± 0.36 kcal/mol
 (Ligand B binds ~245× stronger than A)

Quality checks:

- Overlap matrix: all values should be > 0.03
- dH/dλ curve: smooth, no jumps
- TI vs BAR: should agree within error bars
- Forward vs reverse: hysteresis < 0.5 kcal/mol

pmx generates diagnostic plots automatically: `overlap.png`, `dhdl-TI.png`, `convergence.png`.

11 Case Study: CDK8 Kinase Inhibitors

11.1 The System

- **Target:** CDK8 kinase (cyclin-dependent kinase involved in transcriptional regulation)
- **Ligands:** Small molecule inhibitors with IC₅₀ values spanning 4 orders of magnitude
- **Transformations:** 54 edges, mostly single heavy atom changes (F→Cl, Me→H, ring modifications)

11.2 Protocol

For each ligand pair:

- 21 λ windows (0.00, 0.05, 0.10, ..., 0.95, 1.00)
- 5 ns equilibration + 20 ns production per window
- Both bound (protein-ligand complex) and unbound (ligand in water box) simulations
- Total: $\sim 8.8 \mu\text{s}$ per edge $\times 54$ edges = $\sim 475 \mu\text{s}$ aggregate simulation time

Hardware: Personal RTX 4090 desktop, which ran smoothly.

11.3 Results and Lessons Learned

Convergence metrics:

- Cycle closure: typically < 0.5 kcal/mol
- BAR overlap: > 0.1 for most windows
- Hysteresis: < 0.3 kcal/mol on average

Challenging transformations:

- Charge-changing: required denser λ spacing (30 windows)
- Ring transformations: slower convergence, needed longer equilibration
- Flexible linkers: required 2-3 $\times$ more sampling

Computational efficiency:

- RTX 4090: 3-4 hours per window (25 ns total)
- Total time per edge: 70-90 hours (21 windows $\times$ 2 legs)
- Parallelization: 10-12 edges simultaneously on single GPU

Accuracy: Correlation with experimental $\Delta\Delta G$: Pearson $r = 0.75$, RMSE = 0.9 kcal/mol. This is typical for well-converged FEP and sufficient for rank-ordering compounds in drug discovery.

11.4 Practical Takeaways

Key Insight

The CDK8 benchmark demonstrates that with modern hardware (consumer-grade GPU!) and mature software (pmx/GROMACS), computing dozens of high-quality relative binding free energies is now routine. The fundamental concepts (Zwanzig, overlap, BAR) directly translate to practical convergence criteria. The bottleneck is no longer computational power—it's knowing when your simulations are converged and trustworthy. That's why understanding the theory matters: it tells you what to check.

Figure 4: Statistical Overlap (The Convergence Problem)




Figure 2: Enter Caption

Figure 5: Nonequilibrium Work Distributions (Realistic)




Figure 3: Enter Caption

Figure 6: Perturbation Network Designs




Figure 4: Enter Caption

A Derivation of the Zwanzig Equation

Derivation A.1. *The derivation connects partition functions to probability distributions. Start with the definition of free energy in terms of the partition function:*

$$G_A = -kT \ln Z_A, \quad Z_A = \int e^{-\beta H_A(\mathbf{x})} d\mathbf{x} \quad (38)$$

$$G_B = -kT \ln Z_B, \quad Z_B = \int e^{-\beta H_B(\mathbf{x})} d\mathbf{x} \quad (39)$$

The free energy difference is:

$$\Delta G = G_B - G_A = -kT \ln \frac{Z_B}{Z_A} \quad (40)$$

Now we manipulate Z_B by inserting $e^{-\beta H_A} \cdot e^{\beta H_A} = 1$:

$$Z_B = \int e^{-\beta H_B(\mathbf{x})} d\mathbf{x} \quad (41)$$

$$= \int e^{-\beta H_B(\mathbf{x})} \cdot e^{-\beta H_A(\mathbf{x})} \cdot e^{\beta H_A(\mathbf{x})} d\mathbf{x} \quad (42)$$

$$= \int e^{-\beta(H_B - H_A)(\mathbf{x})} \cdot e^{-\beta H_A(\mathbf{x})} d\mathbf{x} \quad (43)$$

Factor out Z_A :

$$Z_B = Z_A \int e^{-\beta(H_B - H_A)(\mathbf{x})} \cdot \frac{e^{-\beta H_A(\mathbf{x})}}{Z_A} d\mathbf{x} \quad (44)$$

The term $e^{-\beta H_A(\mathbf{x})}/Z_A$ is exactly the Boltzmann probability distribution at state A, so:

$$Z_B = Z_A \left\langle e^{-\beta(H_B - H_A)} \right\rangle_A \quad (45)$$

Taking the logarithm:

$$\ln Z_B = \ln Z_A + \ln \left\langle e^{-\beta(H_B - H_A)} \right\rangle_A \quad (46)$$

Multiply by $-kT$:

$$-kT \ln Z_B = -kT \ln Z_A - kT \ln \left\langle e^{-\beta(H_B - H_A)} \right\rangle_A \quad (47)$$

Therefore:

$$\boxed{\Delta G = G_B - G_A = -kT \ln \left\langle e^{-\beta(H_B - H_A)} \right\rangle_A} \quad (48)$$

B Derivation of Thermodynamic Integration

Starting from the definition of free energy:

$$G(\lambda) = -kT \ln Z(\lambda) \quad (49)$$

Take the derivative with respect to :

$$\frac{dG}{d\lambda} = -kT \frac{d}{d\lambda} \ln Z(\lambda) = -kT \frac{1}{Z(\lambda)} \frac{dZ(\lambda)}{d\lambda} \quad (50)$$

Now compute the derivative of the partition function:

$$\frac{dZ(\lambda)}{d\lambda} = \frac{d}{d\lambda} \int e^{-\beta H(\mathbf{x}, \lambda)} d\mathbf{x} = \int \frac{\partial}{\partial \lambda} e^{-\beta H(\mathbf{x}, \lambda)} d\mathbf{x} \quad (51)$$

Using the chain rule:

$$\frac{\partial}{\partial \lambda} e^{-\beta H(\mathbf{x}, \lambda)} = e^{-\beta H(\mathbf{x}, \lambda)} \cdot \left(-\beta \frac{\partial H}{\partial \lambda} \right) \quad (52)$$

Therefore:

$$\frac{dZ(\lambda)}{d\lambda} = -\beta \int e^{-\beta H(\mathbf{x}, \lambda)} \frac{\partial H}{\partial \lambda} d\mathbf{x} \quad (53)$$

Substituting back:

$$\frac{dG}{d\lambda} = -kT \frac{1}{Z(\lambda)} \left(-\beta \int e^{-\beta H(\mathbf{x}, \lambda)} \frac{\partial H}{\partial \lambda} d\mathbf{x} \right) \quad (54)$$

Since $\beta = 1/(kT)$, the $-kT$ and $-\beta$ cancel:

$$\frac{dG}{d\lambda} = \frac{1}{Z(\lambda)} \int e^{-\beta H(\mathbf{x}, \lambda)} \frac{\partial H}{\partial \lambda} d\mathbf{x} \quad (55)$$

Recognizing that this is an ensemble average:

$$\frac{dG}{d\lambda} = \left\langle \frac{\partial H}{\partial \lambda} \right\rangle_{\lambda} \quad (56)$$

Finally, integrating from $\lambda=0$ to $\lambda=1$:

$$\Delta G = \int_0^1 \frac{dG}{d\lambda} d\lambda = \int_0^1 \left\langle \frac{\partial H}{\partial \lambda} \right\rangle_{\lambda} d\lambda \quad (57)$$

This yields the thermodynamic integration formula:

$$\Delta G = \int_0^1 \left\langle \frac{\partial H}{\partial \lambda} \right\rangle_{\lambda} d\lambda \quad (58)$$

C Derivation of the Jarzynski Equality

This derivation requires careful handling of nonequilibrium statistical mechanics. We give the key steps.

Consider an ensemble of systems, all starting from equilibrium at state A (with distribution $\rho_A(\mathbf{x}) \propto e^{-\beta H_A(\mathbf{x})}$).

For each initial condition $\mathbf{x}_0$, we evolve deterministically (or stochastically) according to the time-dependent Hamiltonian. After time τ , the system is in some final state $\mathbf{x}_{\tau}$.

The work performed along this trajectory is:

$$W[\mathbf{x}(t)] = \int_0^{\tau} \frac{\partial H(\mathbf{x}(t), t)}{\partial t} dt = H(\mathbf{x}_{\tau}, \tau) - H(\mathbf{x}_0, 0) - Q \quad (59)$$

where Q is the heat dissipated (by the first law: $W = \Delta E - Q$).

Now consider the quantity:

$$\langle e^{-\beta W} \rangle = \int \rho_A(\mathbf{x}_0) e^{-\beta W[\mathbf{x}(t)]} d\mathbf{x}_0 \quad (60)$$

The key insight is that the phase space volume element transforms under time evolution. After some calculation (involving the Liouville equation and the Hamiltonian flow), one can show:

$$\int \rho_A(\mathbf{x}_0) e^{-\beta W[\mathbf{x}(t)]} d\mathbf{x}_0 = \frac{Z_B}{Z_A} \quad (61)$$

Therefore:

$$\langle e^{-\beta W} \rangle = \frac{Z_B}{Z_A} = e^{-\beta(G_B - G_A)} = e^{-\beta \Delta G} \quad (62)$$

Taking the logarithm:

$$\Delta G = -kT \ln \langle e^{-\beta W} \rangle \quad (63)$$

The full derivation is subtle and involves measure-theoretic arguments. See Jarzynski’s original 1997 PRL paper for complete details.

D Bennett Acceptance Ratio Derivation

(Optional detailed derivation for interested readers)

The BAR method finds the free energy difference that maximizes the likelihood of observing the sampled work values from both forward and reverse simulations.

Given n_F samples from forward (A→B) with work values $\{W_i^F\}$ and n_R samples from reverse (B→A) with work values $\{W_j^R\}$, the likelihood is:

$$\mathcal{L}(\Delta G) = \prod_i P_F(W_i^F | \Delta G) \prod_j P_R(W_j^R | \Delta G) \quad (64)$$

Taking the logarithm and maximizing (setting $\partial \ln \mathcal{L} / \partial \Delta G = 0$), one arrives at:

$$\sum_{i=1}^{n_F} f(\beta(W_i^F - \Delta G) - C) = \sum_{j=1}^{n_R} f(\beta(-W_j^R - \Delta G) + C) \quad (65)$$

where $f(x) = 1/(1 + e^x)$ is the Fermi function and $C = \ln(n_F/n_R)$.

This is solved iteratively for ΔG .

E Software Comparison Table

Software	License	GPU	Learning Curve	Best For
GROMACS+pmx	Free	Excellent	Medium	Research, large-scale
AMBER+FESetup	Academic free	Good	Medium-High	Versatility
Schrödinger FEP+	Commercial	Excellent	Low	Industry, ease of use
OpenMM+perses	Free	Excellent	High	Method development
NAMD+BFEE	Free	Good	High	Special systems

Table 2: Comparison of major FEP software packages

F Glossary

Alchemical transformation: Non-physical pathway connecting two states, used as a mathematical tool to compute free energy differences.

BAR (Bennett Acceptance Ratio): Optimal estimator for free energy differences given bidirectional sampling.

Cycle closure: Validation check where free energies around a thermodynamic cycle should sum to zero.

Ensemble average: Average of a quantity over all configurations, weighted by Boltzmann probability.

Ergodicity: Assumption that time averages equal ensemble averages (violated in practice by slow dynamics).

FEP (Free Energy Perturbation): Method computing free energy via exponential averaging of energy differences.

Hamiltonian: Total energy function of the system (kinetic + potential).

Hysteresis: Difference between forward and reverse free energy calculations, indicating convergence issues.

Lambda (λ): Alchemical coupling parameter, typically $\lambda \in [0, 1]$.

Overlap: Measure of similarity between probability distributions at adjacent λ windows.

Partition function: $Z = \sum_i e^{-\beta E_i}$, normalization of Boltzmann distribution.

Soft-core potential: Modified interaction that avoids singularities during alchemical transformations.

TI (Thermodynamic Integration): Method computing free energy by integrating $\langle \partial H / \partial \lambda \rangle$.

Work (nonequilibrium): Energy transferred to system during time-dependent perturbation.

References

- [1] Zwanzig, R. W. (1954). High-temperature equation of state by a perturbation method. I. Nonpolar gases. *J. Chem. Phys.* **22**, 1420-1426.
- [2] Kirkwood, J. G. (1935). Statistical mechanics of fluid mixtures. *J. Chem. Phys.* **3**, 300-313.
- [3] Jarzynski, C. (1997). Nonequilibrium equality for free energy differences. *Phys. Rev. Lett.* **78**, 2690-2693.
- [4] Crooks, G. E. (1999). Entropy production fluctuation theorem and the nonequilibrium work relation for free energy differences. *Phys. Rev. E* **60**, 2721-2726.
- [5] Bennett, C. H. (1976). Efficient estimation of free energy differences from Monte Carlo data. *J. Comput. Phys.* **22**, 245-268.
- [6] Shirts, M. R., & Chodera, J. D. (2008). Statistically optimal analysis of samples from multiple equilibrium states. *J. Chem. Phys.* **129**, 124105.
- [7] Klimovich, P. V., Shirts, M. R., & Mobley, D. L. (2015). Guidelines for the analysis of free energy calculations. *J. Comput. Aided Mol. Des.* **29**, 397-411.
- [8] Pohorille, A., Jarzynski, C., & Chipot, C. (2010). Good practices in free-energy calculations. *J. Phys. Chem. B* **114**, 10235-10253.
- [9] Christ, C. D., Mark, A. E., & van Gunsteren, W. F. (2010). Basic ingredients of free energy calculations: A review. *J. Comput. Chem.* **31**, 1569-1582.
- [10] Jarzynski, C. (2011). Equalities and inequalities: Irreversibility and the second law of thermodynamics at the nanoscale. *Annu. Rev. Condens. Matter Phys.* **2**, 329-351.
- [11] Park, S., Khalili-Araghi, F., Tajkhorshid, E., & Schulten, K. (2003). Free energy calculation from steered molecular dynamics simulations using Jarzynski's equality. *J. Chem. Phys.* **119**, 3559-3566.
- [12] Cournia, Z., Allen, B., & Sherman, W. (2017). Relative binding free energy calculations in drug discovery: Recent advances and practical considerations. *J. Chem. Inf. Model.* **57**, 2911-2937.
- [13] Wang, L., Wu, Y., Deng, Y., et al. (2015). Accurate and reliable prediction of relative ligand binding potency in prospective drug discovery by way of a modern free-energy calculation protocol and force field. *J. Am. Chem. Soc.* **137**, 2695-2703.

- [14] Gapsys, V., Michielssens, S., Seeliger, D., & de Groot, B. L. (2015). pmx: Automated protein structure and topology generation for alchemical perturbations. *J. Comput. Chem.* **36**, 348-354.
- [15] Gapsys, V., Pérez-Benito, L., Aldeghi, M., et al. (2020). Large scale relative protein ligand binding affinities using non-equilibrium alchemy. *Chem. Sci.* **11**, 1140-1152.
- [16] Abraham, M. J., Murtola, T., Schulz, R., et al. (2015). GROMACS: High performance molecular simulations through multi-level parallelism from laptops to supercomputers. *SoftwareX* **1-2**, 19-25.
- [17] Beutler, T. C., Mark, A. E., van Schaik, R. C., Gerber, P. R., & van Gunsteren, W. F. (1994). Avoiding singularities and numerical instabilities in free energy calculations based on molecular simulations. *Chem. Phys. Lett.* **222**, 529-539.
- [18] Mobley, D. L., & Klimovich, P. V. (2012). Perspective: Alchemical free energy calculations for drug discovery. *J. Chem. Phys.* **137**, 230901.